


Announcements

Introducing Claude Opus 4.6

Feb 5, 2026



We're upgrading our smartest model.


The new Claude Opus 4.6 improves on its predecessor's coding skills. It plans more carefully, sustains agentic tasks for longer, can operate more reliably in larger codebases, and has better code review and debugging skills to catch its own mistakes. And, in a first for our Opus-class models, Opus 4.6 features a 1M token context window in beta¹.

Opus 4.6 can also apply its improved abilities to a range of everyday work tasks: running financial analyses, doing research, and using and creating documents, spreadsheets, and presentations. Within Cowork, where Claude can multitask autonomously, Opus 4.6 can put all these skills to work on your behalf.

The model’s performance is state-of-the-art on several evaluations. For example, it achieves the highest score on the agentic coding evaluation [Terminal-Bench 2.0](#) and leads all other frontier models on [Humanity’s Last Exam](#), a complex multidisciplinary reasoning test. On [GDPval-AA](#)—an evaluation of performance on economically valuable knowledge work tasks in finance, legal, and other domains²—Opus 4.6 outperforms the industry’s next-best model (OpenAI’s GPT-5.2) by around 144 Elo points,³ and its own predecessor (Claude Opus 4.5) by 190 points. Opus 4.6 also performs better than any other model on [BrowseComp](#), which measures a model’s ability to locate hard-to-find information online.

As we show in our extensive [system card](#), Opus 4.6 also shows an overall safety profile as good as, or better than, any other frontier model in the industry, with low rates of misaligned behavior across safety evaluations.


- Knowledge work
- Agentic search
- Coding
- Reasoning



Opus 4.6 is state-of-the-art on real-world work tasks across several professional domains.

In Claude Code, you can now assemble agent teams to work on tasks together. On the API, Claude can use compaction to summarize its own context and perform longer-running tasks without bumping up against limits. We're also introducing adaptive thinking, where the model can pick up on contextual clues about how much to use its extended thinking, and new effort controls to give developers more control over intelligence, speed, and cost.

We've made substantial upgrades to Claude in Excel, and we're releasing Claude in PowerPoint in a research preview. This makes Claude much more capable for everyday work.



Claude Opus 4.6 is available today on claude.ai, our API, and all major cloud platforms. If you're a developer, use `claude-opus-4-6` via the [Claude API](#). Pricing remains the same at \$5/\$25 per million tokens; for full details, see our [pricing page](#).

We cover the model, our new product updates, our evaluations, and our extensive safety testing in depth below.

First impressions

We build Claude with Claude. Our engineers write code with Claude Code every day, and every new model first gets tested on our own work. With Opus 4.6, we've found that the model brings more focus to the most challenging parts of a task without being told to, moves quickly through the more straightforward parts, handles ambiguous problems with better judgment, and stays productive over longer sessions.

Opus 4.6 often thinks more deeply and more carefully revisits its reasoning before settling on an answer. This produces better results on harder problems, but can add cost and latency on simpler ones. If you're finding that the model is overthinking on a given task, we recommend dialing effort down from its default setting (high) to medium. You can control this easily with the `/effort` [parameter](#).

Here are some of the things our Early Access partners told us about Claude Opus 4.6, including its propensity to work autonomously without hand-holding, its success where previous models failed, and its effect on how teams work:



Notion

“

Claude Opus 4.6 is the strongest model Anthropic has shipped. It takes complicated requests and actually follows through, breaking them into concrete steps, executing, and producing polished work even when the task is ambitious. For Notion users, it feels less like a tool and more like a capable collaborator.

Sarah Sachs
AI Lead, Notion

Evaluating Claude Opus 4.6

Across agentic coding, computer use, tool use, search, and finance, Opus 4.6 is an industry-leading model, often by a wide margin. The table below shows how Claude Opus 4.6 compares to our previous models and to other industry models on a variety of benchmarks.

	Opus 4.6	Opus 4.5	Sonnet 4.5	Gemini 3 Pro	GPT-5.2 (all models)
Agentic terminal coding Terminal-Bench 2.0	65.4%	59.8%	51.0%	56.2% (54.2% self-reported)	64.7% (64.0% self-reported) (Codex CLI)
Agentic coding SWE-bench Verified	80.8%	80.9%	77.2%	76.2%	80.0%
Agentic computer use OSWorld	72.7%	66.3%	61.4%	—	—
Agentic tool use t2-bench	Retail 91.9% Telecom 99.3%	Retail 88.9% Telecom 98.2%	Retail 86.2% Telecom 98.0%	Retail 85.3% Telecom 98.0%	Retail 82.0% Telecom 98.7%
Scaled tool use MCP Atlas	59.5%	62.3%	43.8%	54.1%	60.6%
Agentic search BrowseComp	84.0%	67.8%	43.9%	59.2% (Deep Research)	77.9% (Pro)
Multidisciplinary reasoning Humanity's Last Exam	40.0% without tools 53.0% with tools	30.8% without tools 43.4% with tools	17.7% without tools 33.6% with tools	37.5% without tools 45.8% with tools	36.6% without tools (Pro) 50.0% with tools (Pro)
Agentic financial analysis Finance Agent	60.7%	55.9%	54.2%	44.1%	56.6% (5.1)
Office tasks GDPVal-AA Elo	1606	1416	1277	1195	1462
Novel problem-solving ARC AGI 2	68.8%	37.6%	13.6%	45.1% (Deep Thinking)	54.2% (Pro)
Graduate-level reasoning GPQA Diamond	91.3%	87.0%	83.4%	91.9%	93.2% (Pro)
Visual reasoning MMMU Pro	73.9% without tools 77.3% with tools	70.6% without tools 73.9% with tools	63.4% without tools 68.9% with tools	81.0% without tools — with tools	79.5% without tools 80.4% with tools
Multilingual Q&A MMMLU	91.1%	90.8%	89.5%	91.8%	89.6%

Opus 4.6 is much better at retrieving relevant information from large sets of documents. This extends to long-context tasks, where it holds and tracks information over hundreds of thousands of tokens with less drift, and picks up buried details that even Opus 4.5 would miss.


A common complaint about AI models is “context rot,” where performance degrades as conversations exceed a certain number of tokens. Opus 4.6 performs

markedly better than its predecessors: on the 8-needle 1M variant of MRCR v2—a needle-in-a-haystack benchmark that tests a model’s ability to retrieve information “hidden” in vast amounts of text—Opus 4.6 scores 76%, whereas Sonnet 4.5 scores just 18.5%. This is a qualitative shift in how much context a model can actually use while maintaining peak performance.

All in all, Opus 4.6 is better at finding information across long contexts, better at reasoning after absorbing that information, and has substantially better expert-level reasoning abilities in general.

Long-context retrieval

Long-context reasoning




Opus 4.6 shows significant improvement in long-context retrieval.

Finally, the charts below show how Claude Opus 4.6 performs on a variety of benchmarks that assess its software engineering skills, multilingual coding

ability, long-term coherence, cybersecurity capabilities, and its life sciences knowledge.

- Root cause analysis
- Multilingual coding
- Long-term coherence
- Cybersecurity
- Life sciences




Opus 4.6 excels at diagnosing complex software failures.

A step forward on safety

These intelligence gains do not come at the cost of safety. On our automated behavioral audit, Opus 4.6 showed a low rate of misaligned behaviors such as deception, sycophancy, encouragement of user delusions, and cooperation with misuse. Overall, it is just as well-aligned as its predecessor, Claude Opus 4.5,

which was our most-aligned frontier model to date. Opus 4.6 also shows the lowest rate of over-refusals—where the model fails to answer benign queries—of any recent Claude model.



The overall misaligned behavior score for each recent Claude model on our automated behavioral audit (described in full in the [Claude Opus 4.6 system card](#)).

For Claude Opus 4.6, we ran the most comprehensive set of safety evaluations of any model, applying many different tests for the first time and upgrading several that we’ve used before. We included new evaluations for user wellbeing, more complex tests of the model’s ability to refuse potentially dangerous requests, and updated evaluations of the model’s ability to surreptitiously perform harmful actions. We also experimented with new methods from [interpretability](#), the science of the inner workings of AI models, to begin to understand why the model behaves in certain ways—and, ultimately, to catch problems that standard testing might miss.

A detailed description of all capability and safety evaluations is available in the [Claude Opus 4.6 system card](#).

We've also applied new safeguards in areas where Opus 4.6 shows particular strengths that might be put to dangerous as well as beneficial uses. In particular, since the model shows enhanced cybersecurity abilities, we've developed six new cybersecurity probes—methods of detecting harmful responses—to help us track different forms of potential misuse.

We're also accelerating the *cyberdefensive* uses of the model, using it to help find and patch vulnerabilities in open-source software (as we describe in our new cybersecurity blog post). We think it's critical that cyberdefenders use AI models like Claude to help level the playing field. Cybersecurity moves fast, and we'll be adjusting and updating our safeguards as we learn more about potential threats; in the near future, we may institute real-time intervention to block abuse.

Product and API updates

We've made substantial updates across Claude, Claude Code, and the Claude Developer Platform to let Opus 4.6 perform at its best.

Claude Developer Platform

On the API, we're giving developers better control over model effort and more flexibility for long-running agents. To do so, we're introducing the following features:

- **Adaptive thinking.** Previously, developers only had a binary choice between enabling or disabling extended thinking. Now, with adaptive thinking, Claude can decide when deeper reasoning would be helpful. At the default effort level (high), the model uses extended thinking when useful, but developers can adjust the effort level to make it more or less selective.
- **Effort.** There are now four effort levels to choose from: low, medium, high (default), and max. We encourage developers to experiment with different options to find what works best.
- **Context compaction (beta).** Long-running conversations and agentic tasks often hit the context window. Context compaction automatically summarizes and replaces older context when the conversation approaches a configurable threshold, letting Claude perform longer tasks without hitting limits.
- **1M token context (beta).** Opus 4.6 is our first Opus-class model with 1M token context. Premium pricing applies for prompts exceeding 200k tokens

(\$10/\$37.50 per million input/output tokens), available only on the Claude Developer Platform.

- **128k output tokens.** Opus 4.6 supports outputs of up to 128k tokens, which lets Claude complete larger-output tasks without breaking them into multiple requests.
- **US-only inference.** For workloads that need to run in the United States, US-only inference is available at 1.1× token pricing.

Product updates

Across Claude and Claude Code, we've added features that allow knowledge workers and developers to tackle harder tasks with more of the tools they use every day.

We've introduced agent teams in Claude Code as a research preview. You can now spin up multiple agents that work in parallel as a team and coordinate autonomously—best for tasks that split into independent, read-heavy work like codebase reviews. You can take over any subagent directly using Shift+Up/Down or tmux.

Claude now also works better with the office tools you already use. Claude in Excel handles long-running and harder tasks with improved performance, and can plan before acting, ingest unstructured data and infer the right structure without guidance, and handle multi-step changes in one pass. Pair that with Claude in PowerPoint, and you can first process and structure your data in Excel, then bring it to life visually in PowerPoint. Claude reads your layouts, fonts, and slide masters to stay on brand, whether you're building from a template or generating a full deck from a description. Claude in PowerPoint is now available in research preview for Max, Team, and Enterprise plans.

Footnotes

[1] The 1M token context window is currently available in beta on the Claude Developer Platform only.

[2] Run independently by Artificial Analysis. See here for full methodological details.

[3] This translates into Claude Opus 4.6 obtaining a higher score than GPT-5.2 on this eval approximately 70% of the time (where 50% of the time would have implied parity

in the scores).

- For GPT-5.2 and Gemini 3 Pro models, we compared the best reported model version in the charts and table.
- **Terminal-Bench 2.0:** We report both scores reproduced on our infrastructure and published scores from other labs. All runs used the Terminus-2 harness, except for OpenAI's Codex CLI. All experiments used 1× guaranteed / 3× ceiling resource allocation and 5–15 samples per task across staggered batches. See system card for details.
- **Humanity's Last Exam:** Claude models run “with tools” were run with web search, web fetch, code execution, programmatic tool calling, context compaction triggered at 50k tokens up to 3M total tokens, max reasoning effort, and adaptive thinking enabled. A domain blocklist was used to decontaminate eval results. See system card for more details.
- **SWE-bench Verified:** Our score was averaged over 25 trials. With a prompt modification, we saw a score of 81.42%.
- **MCP Atlas:** Claude Opus 4.6 was run with max effort. When run at high effort, it reached an industry-leading score of 62.7%.
- **BrowseComp:** Claude models were run with web search, web fetch, programmatic tool calling, context compaction triggered at 50k tokens up to 10M total tokens, max reasoning effort, and no thinking enabled. Adding a multi-agent harness increased scores to 86.8%. See system card for more details.
- **ARC AGI 2:** Claude Opus 4.6 was run with max effort and a 120k thinking budget score.
- **CyberGym:** Claude models were run on no thinking, default effort, temperature, and top_p. The model was also given a “think” tool that allowed interleaved thinking for multi-turn evaluations.
- **OpenRCA:** For each failure case in OpenRCA, Claude receives 1 point if all generated root-cause elements match the ground-truth ones, and 0 points if any mismatch is identified. The overall accuracy is the average score across all failure cases. The benchmark was run on the benchmark author's harness, graded using their official methodology, and has been submitted for official verification.

[Feb 23, 2026] Updated reported score for Opus 4.6 for HLE with tools (53.1% to 53.0%). The update was caused by running an improved cheating detection pipeline which flagged 3 additional instances of cheating that our original pipeline had missed.

Related content

Statement from Dario Amodei on our discussions with the Department of War

A statement from our CEO on national security uses of AI.

Read more →

Anthropic acquires Vercept to advance Claude's computer use capabilities

Read more →

Anthropic's Responsible Scaling Policy: Version 3.0

Read more →