


Product

Introducing Claude Sonnet 4.6

Feb 17, 2026



Claude Sonnet 4.6 is our most capable Sonnet model yet. It's a full upgrade of the model's skills across coding, computer use, long-context reasoning, agent planning, knowledge work, and design. Sonnet 4.6 also features a 1M token context window in beta.

For those on our [Free and Pro plans](#), Claude Sonnet 4.6 is now the default model in [Claude.ai](#) and [Claude Cowork](#). [Pricing](#) remains the same as Sonnet 4.5, starting at \$3/\$15 per million tokens.

Sonnet 4.6 brings much-improved coding skills to more of our users. Improvements in consistency, instruction following, and more have made developers with early access prefer Sonnet 4.6 to its predecessor by a wide margin. They often even prefer it to our smartest model from November 2025, Claude Opus 4.5.

Performance that would have previously required reaching for an Opus-class model—including on real-world, economically valuable [office tasks](#)—is now available with Sonnet 4.6. The model also shows a major improvement in computer use skills compared to prior Sonnet models.

As with every new Claude model, we've run [extensive safety evaluations](#) of Sonnet 4.6, which overall showed it to be as safe as, or safer than, our other recent Claude models. Our safety researchers concluded that Sonnet 4.6 has “a broadly warm, honest, prosocial, and at times funny character, very strong safety behaviors, and no signs of major concerns around high-stakes forms of misalignment.”


Computer use

Almost every organization has software it can't easily automate: specialized systems and tools built before modern interfaces like APIs existed. To have AI use such software, users would previously have had to build bespoke connectors. But a model that can use a computer the way a person does changes that equation.

In October 2024, we were the first to introduce a general-purpose computer-using model. At the time, we wrote that it was “still experimental—at times cumbersome and error-prone,” but we expected rapid improvement. OSWorld, the standard benchmark for AI computer use, shows how far our models have come. It presents hundreds of tasks across real software (Chrome, LibreOffice, VS Code, and more) running on a simulated computer. There are no special APIs or purpose-built connectors; the model sees the computer and interacts with it in much the same way a person would: clicking a (virtual) mouse and typing on a (virtual) keyboard.

Across sixteen months, our Sonnet models have made steady gains on OSWorld. The improvements can also be seen beyond benchmarks: early Sonnet 4.6 users are seeing human-level capability in tasks like navigating a complex spreadsheet or filling out a multi-step web form, before pulling it all together across multiple browser tabs.

The model certainly still lags behind the most skilled humans at using computers. But the rate of progress is remarkable nonetheless. It means that computer use is much more useful for a range of work tasks—and that substantially more capable models are within reach.



Scores prior to Claude Sonnet 4.5 were measured on the original OSWorld; scores from Sonnet 4.5 onward use OSWorld-Verified. OSWorld-Verified (released July 2025) is an in-place upgrade of the original OSWorld benchmark, with updates to task quality, evaluation grading, and infrastructure.

At the same time, computer use poses risks: malicious actors can attempt to hijack the model by hiding instructions on websites in what’s known as a prompt injection attack. We’ve been working to improve our models’ resistance to prompt injections—our safety evaluations show that Sonnet 4.6 is a major improvement compared to its predecessor, Sonnet 4.5, and performs similarly to Opus 4.6. You can find out more about how to mitigate prompt injections and other safety concerns in our API docs.

Evaluating Claude Sonnet 4.6

Beyond computer use, Claude Sonnet 4.6 has improved on benchmarks across the board. It approaches Opus-level intelligence at a price point that makes it more practical for far more tasks. You can find a full discussion of Sonnet 4.6’s capabilities and its safety-related behaviors in our system card; a summary and comparison to other recent models is below.

	Sonnet 4.6	Sonnet 4.5	Opus 4.6	Opus 4.5	Gemini 3 Pro	GPT-5.2 (all models)
Agentic terminal coding Terminal-Bench 2.0	59.1%	51.0%	65.4%	59.8%	56.2% (54.2% self-reported)	64.7% (64.0% self-reported) (Codex CLI)
Agentic coding SWE-bench Verified	79.6%	77.2%	80.8%	80.9%	78.0% (Flash)	80.0%
Agentic computer use OSWorld-Verified	72.5%	61.4%	72.7%	66.3%	—	38.2%
Agentic tool use t2-bench	Retail 91.7%	Retail 86.2%	Retail 91.9%	Retail 88.9%	Retail 85.3%	Retail 82.0%
	Telecom 97.9%	Telecom 98.0%	Telecom 99.3%	Telecom 98.2%	Telecom 98.0%	Telecom 98.7%
Scaled tool use MCP-Atlas	61.3%	43.8%	59.5%	62.3%	54.1%	60.6%
Agentic search BrowseComp	74.7%	43.9%	84.0%	67.8%	59.2% (Deep Research)	77.9% (Pro)
Multidisciplinary reasoning Humanity's Last Exam (HLE)	33.2% without tools	17.7% without tools	40.0% without tools	30.8% without tools	37.5% without tools	36.6% without tools (Pro)
	49.0% with tools	33.6% with tools	53.0% with tools	43.4% with tools	45.8% with tools	50.0% with tools (Pro)
Agentic financial analysis Finance Agent v1.1	63.3%	54.5%	60.1%	58.8%	55.2%	59.0%
Office tasks GDPval-AA Elo	1633	1276	1606	1416	1201	1462
Novel problem-solving ARC-AGI-2	58.3%	13.6%	68.8%	37.6%	31.1%	54.2% (Pro)
Graduate-level reasoning GPQA Diamond	89.9%	83.4%	91.3%	87.0%	91.9%	93.2% (Pro)
Visual reasoning MMMU-Pro	74.5% without tools	63.4% without tools	73.9% without tools	70.6% without tools	81.0% without tools	79.5% without tools
	75.6% with tools	68.9% with tools	77.3% with tools	73.9% with tools	— with tools	80.4% with tools
Multilingual Q&A MMMLU	89.3%	89.5%	91.1%	90.8%	91.8%	89.6%

In Claude Code, our early testing found that users preferred Sonnet 4.6 over Sonnet 4.5 roughly 70% of the time. Users reported that it more effectively read the context before modifying code and consolidated shared logic rather than duplicating it. This made it less frustrating to use over long sessions than earlier models.


Users even preferred Sonnet 4.6 to Opus 4.5, our frontier model from November, 59% of the time. They rated Sonnet 4.6 as significantly less prone to overengineering and “laziness,” and meaningfully better at instruction following. They reported fewer false claims of success, fewer hallucinations, and more consistent follow-through on multi-step tasks.

Sonnet 4.6’s 1M token context window is enough to hold entire codebases, lengthy contracts, or dozens of research papers in a single request. More importantly, Sonnet 4.6 *reasons effectively* across all that context. This can make it much better at long-horizon planning. We saw this particularly clearly in the [Vending-Bench Arena](#) evaluation, which tests how well a model can run a (simulated) business over time—and which includes an element of competition, with different AI models facing off against each other to make the biggest profits.

Sonnet 4.6 developed an interesting new strategy: it invested heavily in capacity for the first ten simulated months, spending significantly more than its competitors, and then pivoted sharply to focus on profitability in the final stretch. The timing of this pivot helped it finish well ahead of the competition.

Money balance over time

Vending-Bench Arena Sonnet 4.6 vs. Sonnet 4.5



Sonnet 4.6 outperforms Sonnet 4.5 on Vending-Bench Arena by investing in capacity early, then pivoting to profitability in the final stretch.

Early customers also reported broad improvements, with frontend code and financial analysis standing out. Customers independently described visual outputs from Sonnet 4.6 as notably more polished, with better layouts, animations, and design sensibility than those from previous models. Customers also needed fewer rounds of iteration to reach production-quality results.

“

Claude Sonnet 4.6 matches Opus 4.6 performance on OfficeQA, which measures how well a model can read enterprise documents (charts, PDFs, tables), pull the right facts, and reason from those facts. It's a meaningful upgrade for document comprehension workloads.

Hanlin Tang
CTO of Neural Networks, Databricks

“

The performance-to-cost ratio of Claude Sonnet 4.6 is extraordinary—it's hard to overstate how fast Claude models have been evolving in recent months. Sonnet 4.6 outperforms on our orchestration evals, handles our most complex agentic workloads, and keeps improving the higher you push the effort settings.

Michele Catasta
President, Replit

Product updates

On the Claude Developer Platform, Sonnet 4.6 supports both [adaptive thinking](#) and extended thinking, as well as [context compaction](#) in beta, which automatically summarizes older context as conversations approach limits, increasing effective context length.

On our API, Claude's [web search](#) and [fetch](#) tools now automatically write and execute code to [filter and process search results](#), keeping only relevant content in context—improving both response quality and token efficiency. Additionally,

[code execution](#), [memory](#), [programmatic tool calling](#), [tool search](#), and [tool use examples](#) are now generally available.

Sonnet 4.6 offers strong performance at any thinking effort, even with extended thinking off. As part of your migration from Sonnet 4.5, we recommend exploring across the spectrum to find the ideal balance of speed and reliable performance, depending on what you're building.

We find that Opus 4.6 remains the strongest option for tasks that demand the deepest reasoning, such as codebase refactoring, coordinating multiple agents in a workflow, and problems where getting it *just right* is paramount.

For [Claude in Excel](#) users, our add-in now supports MCP connectors, letting Claude work with the other tools you use day-to-day, like S&P Global, LSEG, Daloopa, PitchBook, Moody's, and FactSet. You can ask Claude to pull in context from outside your spreadsheet without ever leaving Excel. If you've already set up MCP connectors in Claude.ai, those same connections will work in Excel automatically. This is available on Pro, Max, Team, and Enterprise plans.

How to use Claude Sonnet 4.6

Claude Sonnet 4.6 is available now on all [Claude plans](#), [Claude Cowork](#), [Claude Code](#), our API, and all major cloud platforms. We've also upgraded our free tier to Sonnet 4.6 by default—it now includes file creation, connectors, skills, and compaction.

If you're a developer, you can get started quickly by using `claude-sonnet-4-6` via the [Claude API](#).

Footnotes

- For GPT-5.2 and Gemini 3 Pro, we compared against the best reported model version available via API in the charts and table.
- OSWorld:** OSWorld tests a specific set of computer tasks in a controlled environment. It's one of the best measures we have, but not a complete picture of real-world computer use. Real-world computer use is often messier and more ambiguous, and it carries higher stakes for errors. No benchmark fully captures that yet.
- Terminal-Bench 2.0:** We report both scores reproduced on our infrastructure and published scores from other labs. All runs used the Terminus-2 harness, except for OpenAI's Codex CLI. All experiments used 1× guaranteed/3× ceiling resource allocation and 5–15 samples per task across staggered batches. The Sonnet 4.6 score reported is with thinking turned off.
- SWE-bench Verified:** Our score was averaged over 10 trials. With a prompt modification, we saw a score of 80.2%.
- Humanity's Last Exam:** Claude models run “with tools” were run with web search, web fetch, code execution, programmatic tool calling, context compaction triggered at 50k tokens up to 3M total tokens, max reasoning effort, and adaptive thinking enabled. A domain blocklist was used to decontaminate eval results.
- BrowseComp:** Claude models were run with web search, web fetch, programmatic tool calling, context compaction triggered at 50k tokens up to 10M total tokens, max reasoning effort, and no thinking enabled.
- ARC-AGI-2:** Claude Sonnet 4.6 was run with max and high effort and a 120k thinking budget score. The score shown reflects max effort; with high effort, we achieve a score of 60.4%.
- MMMU-Pro:** We made two small updates to our MMMU-Pro implementation that have affected the score: 1) our previous implementation contained the prefix “Let's think step-by-step,” which we have removed, and 2) we previously graded this multiple-choice eval by looking at on-policy token probabilities of the multiple-choice options; we now grade it using a separate model (Claude Sonnet 4).



Related content

Statement from Dario Amodei on our discussions with the Department of War

A statement from our CEO on national security uses of AI.

[Read more](#) →

Anthropic acquires Vercept to advance Claude's computer use capabilities

[Read more](#) →

Anthropic's Responsible Scaling Policy: Version 3.0

[Read more](#) →



Products

Claude
Claude Code
Claude Code Enterprise
Cowork
Claude in Chrome
Claude in Excel
Claude in PowerPoint
Claude in Slack
Skills
Max plan
Team plan
Enterprise plan
Download app
Pricing
Log in to Claude

Models

Opus
Sonnet
Haiku

Solutions

AI agents
Code modernization
Coding
Customer support
Education
Financial services
Government
Healthcare
Life sciences
Nonprofits

Claude Developer Platform

Overview
Developer docs
Pricing
Regional compliance
Amazon Bedrock
Google Cloud's Vertex AI
Console login

Learn

Blog
Claude partner network
Connectors
Courses
Customer stories
Engineering at Anthropic
Events
Plugins
Powered by Claude
Service partners
Startups program
Tutorials
Use cases

Company

Anthropic
Careers
Economic Futures
Research
News
Claude's Constitution
Responsible Scaling Policy
Security and compliance
Transparency

Help and security

Availability
Status
Support center

Terms and policies
Privacy policy
Consumer health data privacy policy
Responsible disclosure policy
Terms of service: Commercial
Terms of service: Consumer
Usage policy